C1_W1_Lab_3_siamese-network

December 23, 2020

1 Ungraded Lab: Implement a Siamese network

This lab will go through creating and training a multi-input model. You will build a basic Siamese Network to find the similarity or dissimilarity between items of clothing. For Week 1, you will just focus on constructing the network. You will revisit this lab in Week 2 when we talk about custom loss functions.

1.1 Imports

```
[1]: try:
     # %tensorflow_version only exists in Colab.
       %tensorflow_version 2.x
     except Exception:
       pass

     import tensorflow as tf
     from tensorflow.keras.models import Model
     from tensorflow.keras.layers import Input, Flatten, Dense, Dropout, Lambda
     from tensorflow.keras.optimizers import RMSprop
     from tensorflow.keras.datasets import fashion_mnist
     from tensorflow.python.keras.utils.vis_utils import plot_model
     from tensorflow.keras import backend as K

     import numpy as np
     import matplotlib.pyplot as plt
     from PIL import Image, ImageFont, ImageDraw
     import random
```

1.2 Prepare the Dataset

First define a few utilities for preparing and visualizing your dataset.

```
[2]: def create_pairs(x, digit_indices):
         '''Positive and negative pair creation.
         Alternates between positive and negative pairs.
```

```
    '''
    pairs = []
    labels = []
    n = min([len(digit_indices[d]) for d in range(10)]) - 1

    for d in range(10):
        for i in range(n):
            z1, z2 = digit_indices[d][i], digit_indices[d][i + 1]
            pairs += [[x[z1], x[z2]]]
            inc = random.randrange(1, 10)
            dn = (d + inc) % 10
            z1, z2 = digit_indices[d][i], digit_indices[dn][i]
            pairs += [[x[z1], x[z2]]]
            labels += [1, 0]

    return np.array(pairs), np.array(labels)


def create_pairs_on_set(images, labels):

    digit_indices = [np.where(labels == i)[0] for i in range(10)]
    pairs, y = create_pairs(images, digit_indices)
    y = y.astype('float32')

    return pairs, y


def show_image(image):
    plt.figure()
    plt.imshow(image)
    plt.colorbar()
    plt.grid(False)
    plt.show()
```

You can now download and prepare our train and test sets. You will also create pairs of images that will go into the multi-input model.

```
[3]: # load the dataset
     (train_images, train_labels), (test_images, test_labels) = fashion_mnist.
      →load_data()

     # prepare train and test sets
     train_images = train_images.astype('float32')
     test_images = test_images.astype('float32')

     # normalize values
     train_images = train_images / 255.0
```

```
test_images = test_images / 255.0

# create pairs on train and test sets
tr_pairs, tr_y = create_pairs_on_set(train_images, train_labels)
ts_pairs, ts_y = create_pairs_on_set(test_images, test_labels)
```

You can see a sample pair of images below.

```
[4]: # array index
     this_pair = 8

     # show images at this index
     show_image(ts_pairs[this_pair][0])
     show_image(ts_pairs[this_pair][1])

     # print the label for this pair
     print(tr_y[this_pair])
```

1.0

```
[5]: # print other pairs

show_image(tr_pairs[:,0][0])
show_image(tr_pairs[:,0][1])

show_image(tr_pairs[:,1][0])
show_image(tr_pairs[:,1][1])
```

1.3 Build the Model

Next, you'll define some utilities for building our model.

```
[6]:  def initialize_base_network():
          input = Input(shape=(28,28,), name="base_input")
          x = Flatten(name="flatten_input")(input)
          x = Dense(128, activation='relu', name="first_base_dense")(x)
          x = Dropout(0.1, name="first_dropout")(x)
          x = Dense(128, activation='relu', name="second_base_dense")(x)
          x = Dropout(0.1, name="second_dropout")(x)
          x = Dense(128, activation='relu', name="third_base_dense")(x)

          return Model(inputs=input, outputs=x)


      def euclidean_distance(vects):
          x, y = vects
          sum_square = K.sum(K.square(x - y), axis=1, keepdims=True)
          return K.sqrt(K.maximum(sum_square, K.epsilon()))


      def eucl_dist_output_shape(shapes):
          shape1, shape2 = shapes
          return (shape1[0], 1)
```

Let's see how our base network looks. This is where the two inputs will pass through to generate an output vector.

```
[7]:  base_network = initialize_base_network()
      plot_model(base_network, show_shapes=True, show_layer_names=True,␣
       ↪to_file='base-model.png')
```

[7]:

base_input: InputLayer	input:	[(?, 28, 28)]
	output:	[(?, 28, 28)]

flatten_input: Flatten	input:	(?, 28, 28)
	output:	(?, 784)

first_base_dense: Dense	input:	(?, 784)
	output:	(?, 128)

first_dropout: Dropout	input:	(?, 128)
	output:	(?, 128)

second_base_dense: Dense	input:	(?, 128)
	output:	(?, 128)

second_dropout: Dropout	input:	(?, 128)
	output:	(?, 128)

third_base_dense: Dense	input:	(?, 128)
	output:	(?, 128)
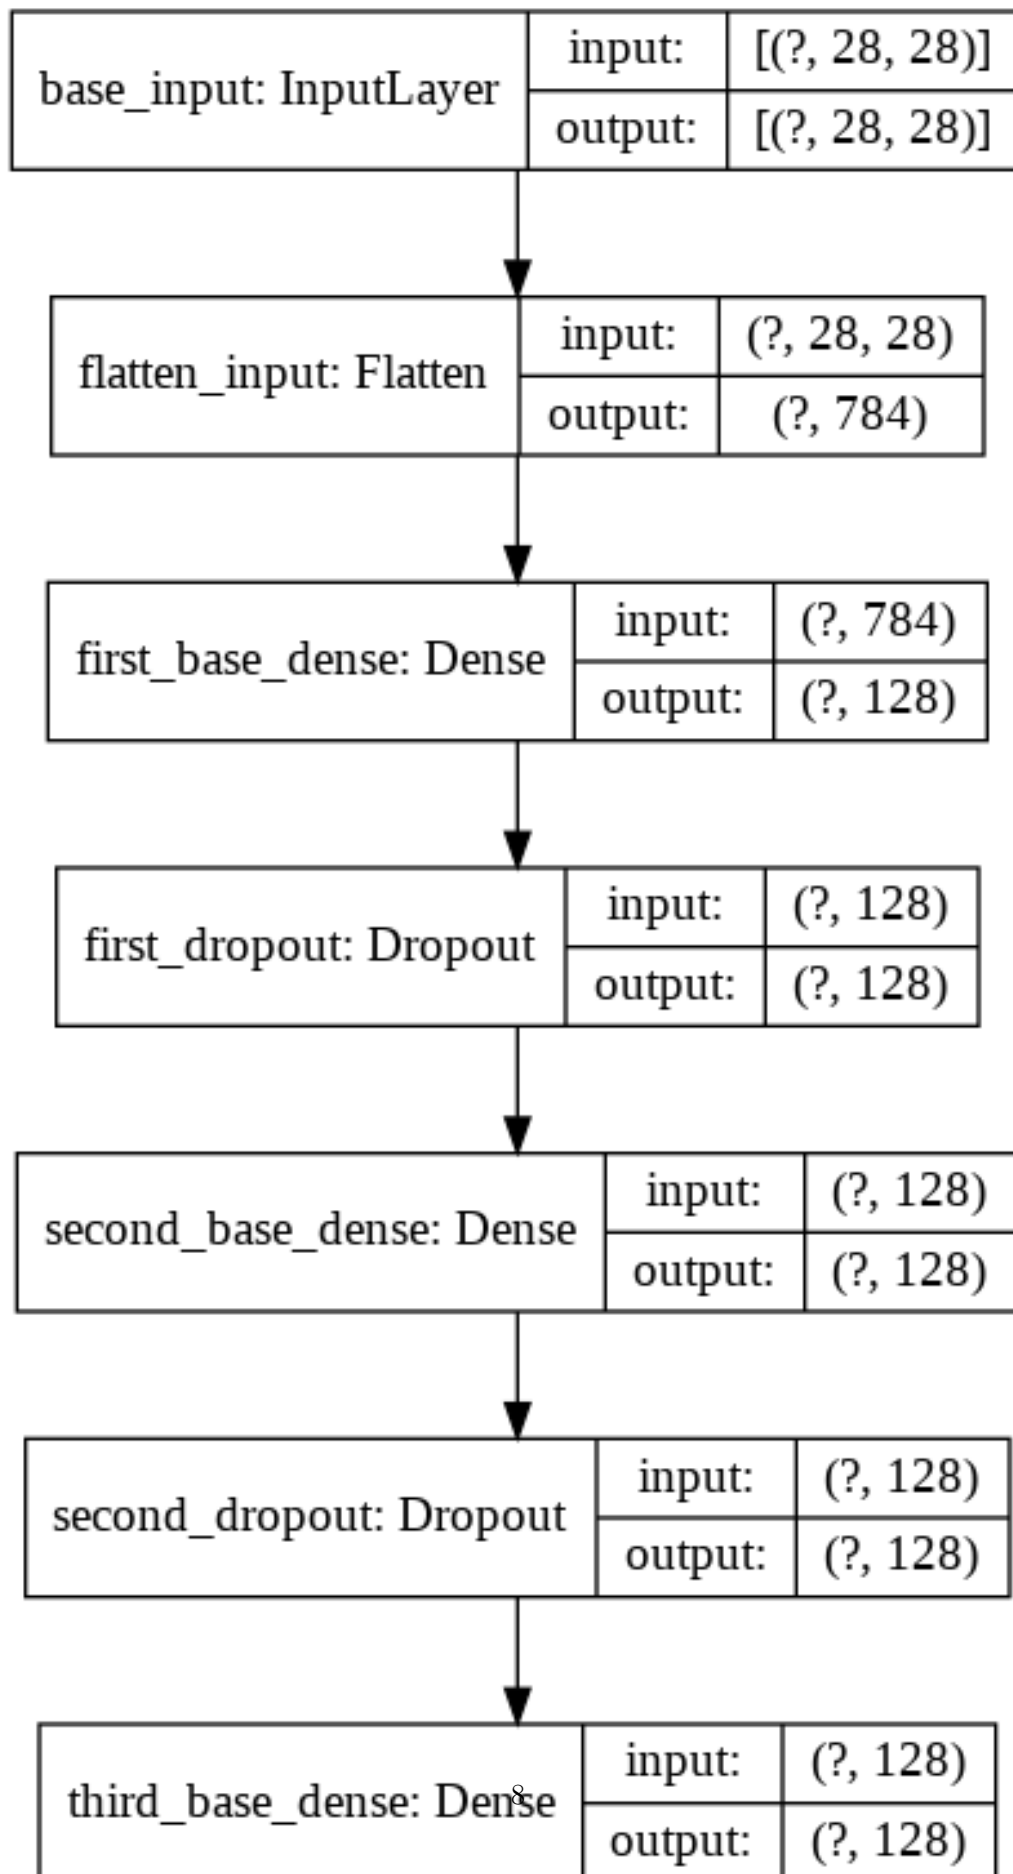
Let's now build the Siamese network. The plot will show two inputs going to the base network.
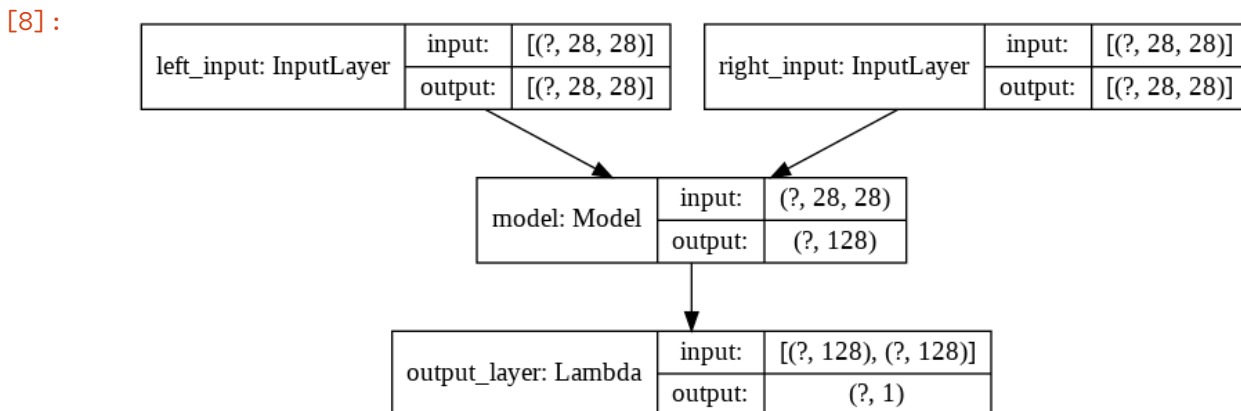
```
[8]: # create the left input and point to the base network
     input_a = Input(shape=(28,28,), name="left_input")
     vect_output_a = base_network(input_a)

     # create the right input and point to the base network
     input_b = Input(shape=(28,28,), name="right_input")
     vect_output_b = base_network(input_b)

     # measure the similarity of the two vector outputs
     output = Lambda(euclidean_distance, name="output_layer",␣
      ↪output_shape=eucl_dist_output_shape)([vect_output_a, vect_output_b])

     # specify the inputs and output of the model
     model = Model([input_a, input_b], output)

     # plot model graph
     plot_model(model, show_shapes=True, show_layer_names=True, to_file='outer-model.
      ↪png')
```

[8]:

left_input: InputLayer	input:	[(?, 28, 28)]
	output:	[(?, 28, 28)]

right_input: InputLayer	input:	[(?, 28, 28)]
	output:	[(?, 28, 28)]

model: Model	input:	(?, 28, 28)
	output:	(?, 128)

output_layer: Lambda	input:	[(?, 128), (?, 128)]
	output:	(?, 1)

1.4 Train the Model

You can now define the custom loss for our network and start training. Don't worry about why it's written as a nested function just yet. You will revisit this in Week 2.

```
[9]: def contrastive_loss_with_margin(margin):
         def contrastive_loss(y_true, y_pred):
```

```
        '''Contrastive loss from Hadsell-et-al.'06
        http://yann.lecun.com/exdb/publis/pdf/hadsell-chopra-lecun-06.pdf
        '''
        square_pred = K.square(y_pred)
        margin_square = K.square(K.maximum(margin - y_pred, 0))
        return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
    return contrastive_loss
```

```
[11]: rms = RMSprop()
      model.compile(loss=contrastive_loss_with_margin(margin=1), optimizer=rms)
      history = model.fit([tr_pairs[:,0], tr_pairs[:,1]], tr_y, epochs=4,␣
       ↪batch_size=128, validation_data=([ts_pairs[:,0], ts_pairs[:,1]], ts_y))
```

```
Train on 119980 samples, validate on 19980 samples
Epoch 1/4
119980/119980 [==============================] - 14s 114us/sample - loss: 0.0691
- val_loss: 0.0709
Epoch 2/4
119980/119980 [==============================] - 8s 67us/sample - loss: 0.0649 -
val_loss: 0.0663
Epoch 3/4
119980/119980 [==============================] - 8s 66us/sample - loss: 0.0623 -
val_loss: 0.0664
Epoch 4/4
119980/119980 [==============================] - 8s 65us/sample - loss: 0.0609 -
val_loss: 0.0643
```

1.5 Model Evaluation

As usual, you can evaluate our model by computing the accuracy and observing the metrics during training.

```
[12]: def compute_accuracy(y_true, y_pred):
        '''Compute classification accuracy with a fixed threshold on distances.
        '''
        pred = y_pred.ravel() > 0.5
        return np.mean(pred == y_true)
```

```
[13]: loss = model.evaluate(x=[ts_pairs[:,0],ts_pairs[:,1]], y=ts_y)

      y_pred_train = model.predict([tr_pairs[:,0], tr_pairs[:,1]])
      train_accuracy = compute_accuracy(tr_y, y_pred_train)

      y_pred_test = model.predict([ts_pairs[:,0], ts_pairs[:,1]])
      test_accuracy = compute_accuracy(ts_y, y_pred_test)
```

```
print("Loss = {}, Train Accuracy = {} Test Accuracy = {}".format(loss,␣
  ↪train_accuracy, test_accuracy))
```

```
19980/19980 [==============================] - 1s 36us/sample - loss: 0.0643
Loss = 0.0642609304149968, Train Accuracy = 0.06932822137022837 Test Accuracy =
0.08418418418418419
```

[14]:
```
def plot_metrics(metric_name, title, ylim=5):
    plt.title(title)
    plt.ylim(0,ylim)
    plt.plot(history.history[metric_name],color='blue',label=metric_name)
    plt.plot(history.history['val_' + metric_name],color='green',label='val_' +␣
  ↪metric_name)


plot_metrics(metric_name='loss', title="Loss", ylim=0.2)
```


[15]:
```
# Matplotlib config
def visualize_images():
    plt.rc('image', cmap='gray_r')
    plt.rc('grid', linewidth=0)
    plt.rc('xtick', top=False, bottom=False, labelsize='large')
    plt.rc('ytick', left=False, right=False, labelsize='large')
    plt.rc('axes', facecolor='F8F8F8', titlesize="large", edgecolor='white')
```

11

```
    plt.rc('text', color='a8151a')
    plt.rc('figure', facecolor='F0F0F0')# Matplotlib fonts



# utility to display a row of digits with their predictions
def display_images(left, right, predictions, labels, title, n):
    plt.figure(figsize=(17,3))
    plt.title(title)
    plt.yticks([])
    plt.xticks([])
    plt.grid(None)
    left = np.reshape(left, [n, 28, 28])
    left = np.swapaxes(left, 0, 1)
    left = np.reshape(left, [28, 28*n])
    plt.imshow(left)
    plt.figure(figsize=(17,3))
    plt.yticks([])
    plt.xticks([28*x+14 for x in range(n)], predictions)
    for i,t in enumerate(plt.gca().xaxis.get_ticklabels()):
        if predictions[i] > 0.5: t.set_color('red') # bad predictions in red
    plt.grid(None)
    right = np.reshape(right, [n, 28, 28])
    right = np.swapaxes(right, 0, 1)
    right = np.reshape(right, [28, 28*n])
    plt.imshow(right)
```
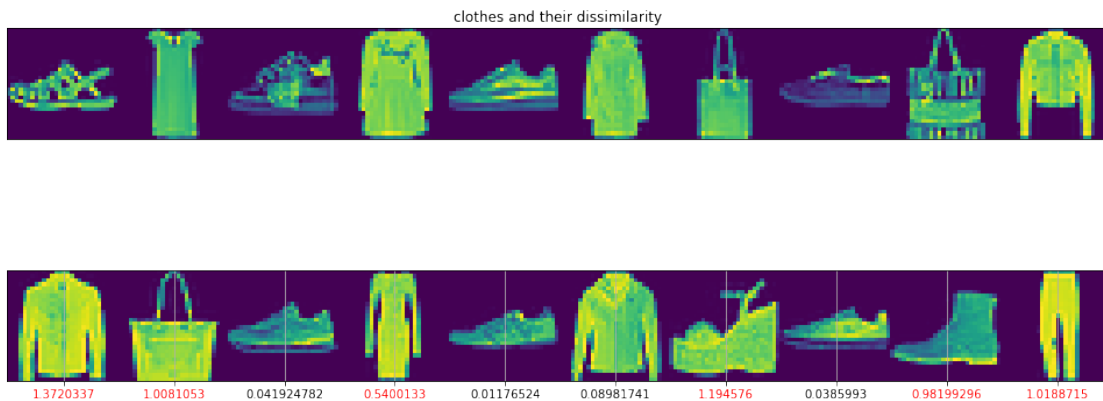
You can see sample results for 10 pairs of items below.

```
[16]: y_pred_train = np.squeeze(y_pred_train)
      indexes = np.random.choice(len(y_pred_train), size=10)
      display_images(tr_pairs[:, 0][indexes], tr_pairs[:, 1][indexes],␣
       →y_pred_train[indexes], tr_y[indexes], "clothes and their dissimilarity", 10)
```


clothes and their dissimilarity

1.3720337   1.0081053   0.041924782   0.5400133   0.01176524   0.08981741   1.194576   0.0385993   0.98199296   1.0188715

```

```